



<https://www.docker.io/>

What is it?

- Docker is an open-source project to easily create lightweight, portable, self-sufficient containers from any application
- Makes fast running, portable **containers**

History

- Dotcloud, Inc creates PaaS service
- January 2013, work starts on docker internally
- March 2013, first public release
- Statistics:
 - 16853 stars on github
 - 3343 forks
 - 684 contributors
- Massive community interest
- Created by Solomon Hykes (French engineer ;)

**EVER TRIED.
EVER FAILED.
NO MATTER.
TRY AGAIN.
FAIL AGAIN.
FAIL BETTER.**

Samuel Beckett (1906-1989)

Why this hype?

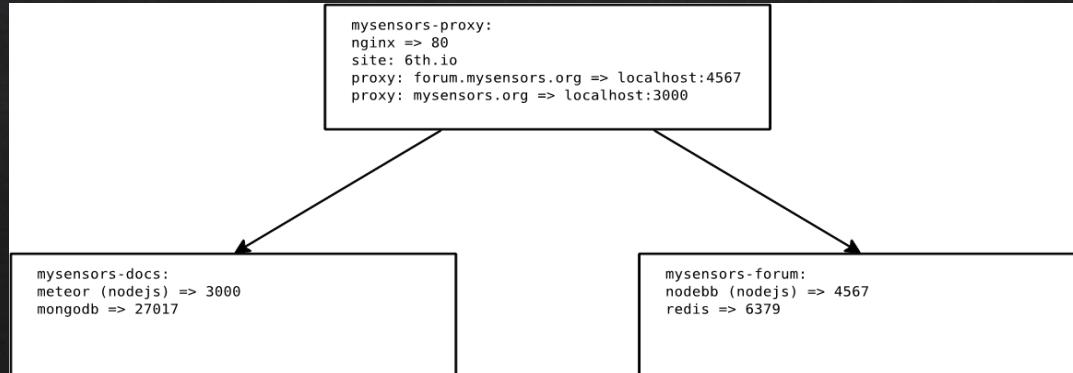
- Solves an important problem
- Easy to use
- Efficient



Why docker?

Problem

- .An application have many dependencies that needs to coexist
- .Deployments with different demands:
 - Development/CI
 - Test
 - Production



Application Deployment got complex

Application Stack

- Basic OS
- JVM
- Static web server
- Front-end platform
- Database layer
- Application code

X

Deployment environments

- Development VM
- QA Server
- Customer Data Center
- Public Cloud
- Contributors Laptop
- Production Servers
- Production Clusters

Problems

- Installing software
- Software versioning
- Configuration
- Testing



What is docker?



Analogy

.Transporting of goods before 1960



Solution: Shipping container



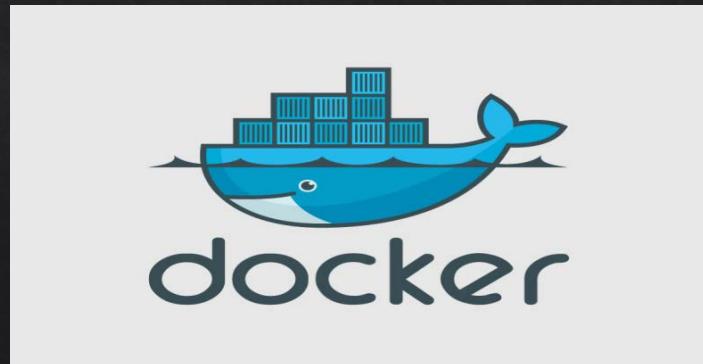
- Separation of concerns
 - User cares about packing the inside
 - Shipper cares about moving the container
- Standardized interface



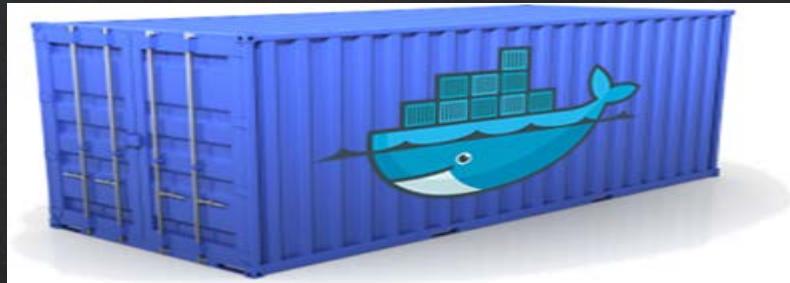
What is docker?



Solution
.Containers



Docker containers



Standardized interface for software container

Developer concerns

- .Code
- .Libraries
- .Services
- .Configuration
- .Data

All servers look the same

Ops concerns

- .Moving containers
- .Starting/Stopping containers
- .Logging
- .Monitoring
- .Network configuration

All containers look the same

Isolation

What is a Container?

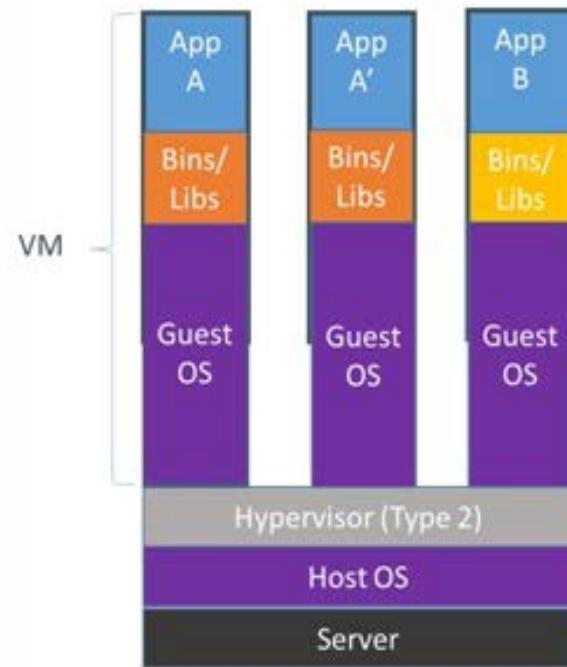
- **Operating system–level virtualization**
- Operating system–level virtualization is a server virtualization method where the kernel of an operating system allows for multiple isolated user-space instances, instead of just one
- Differs from traditional virtual machines

Containers vs. VMs

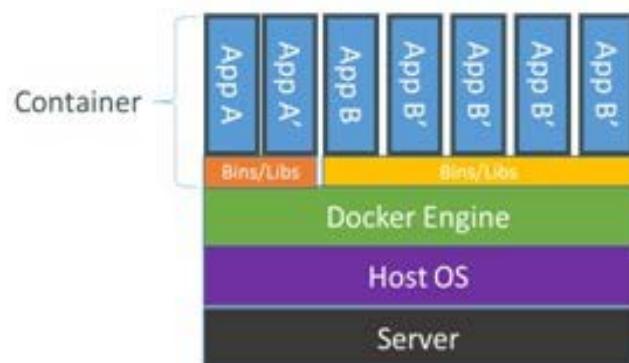
- Container
 - Shares the host OS and kernel
 - Zero boot time
 - Cannot run any OS (strictly Linux for Docker)
 - Little to no start-up or performance penalty. Technically native.
- Virtualization
 - Full OSes on top of a host OS via hypervisor
 - Full software stack
 - Each VM has its own kernel
 - Full boot process for each VM (slow)
 - Can run any OS (Windows and BSD included)

Containers vs. VMs (Pretty Picture)

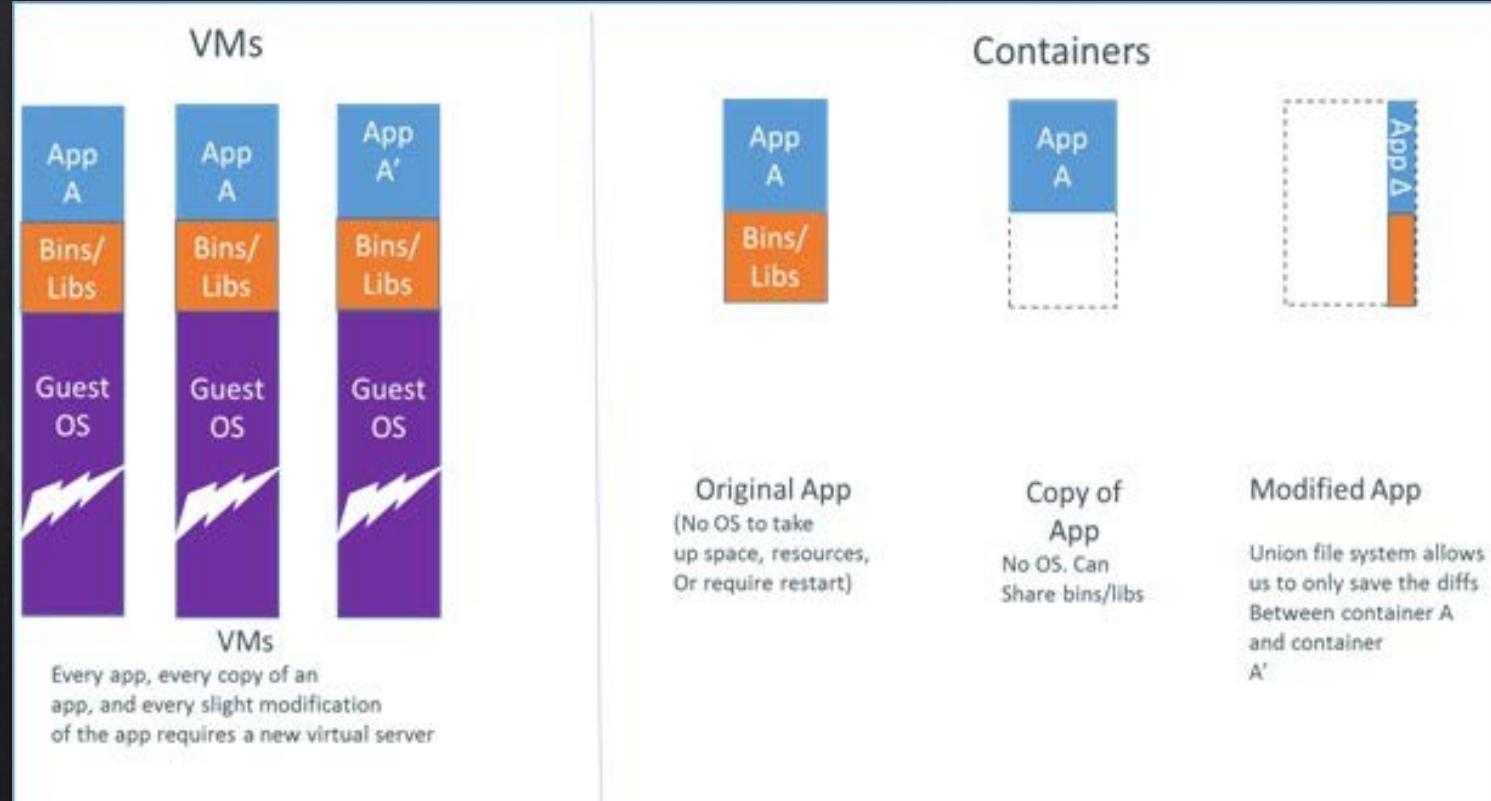
Containers vs. VMs



Containers are isolated,
but share OS and, where
appropriate, bins/libraries

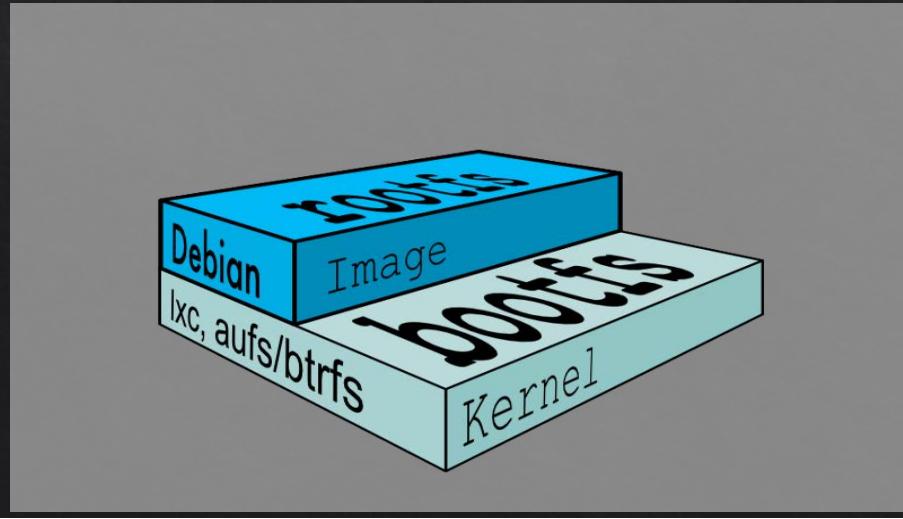


Containers vs. VMs (Another Pretty Picture)



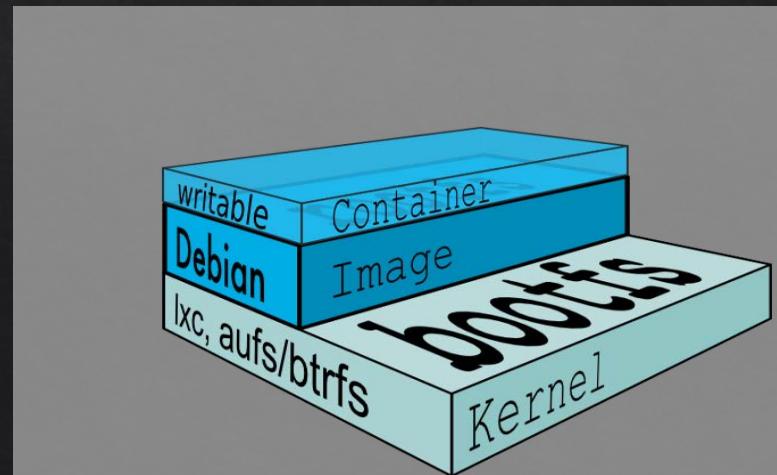
Docker glossary

- **Image**
 - Read-only template for a container
 - Includes all files required for application to run
 - Has additional metadata
 - Exposed network ports
 - Binary to start



Docker glossary

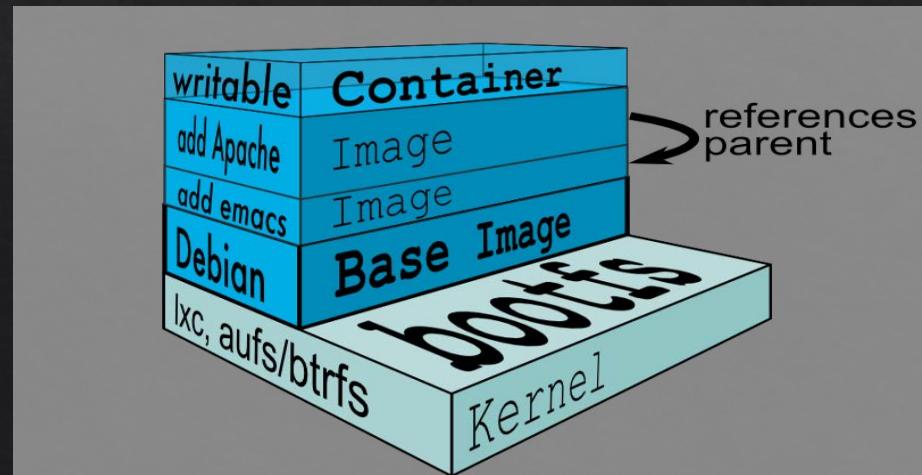
- Container
 - Running processes
 - Based on a particular image
 - Typically a single process
 - Isolated from host system
 - Cheap
 - Can write to filesystem
 - Commit creates new Image

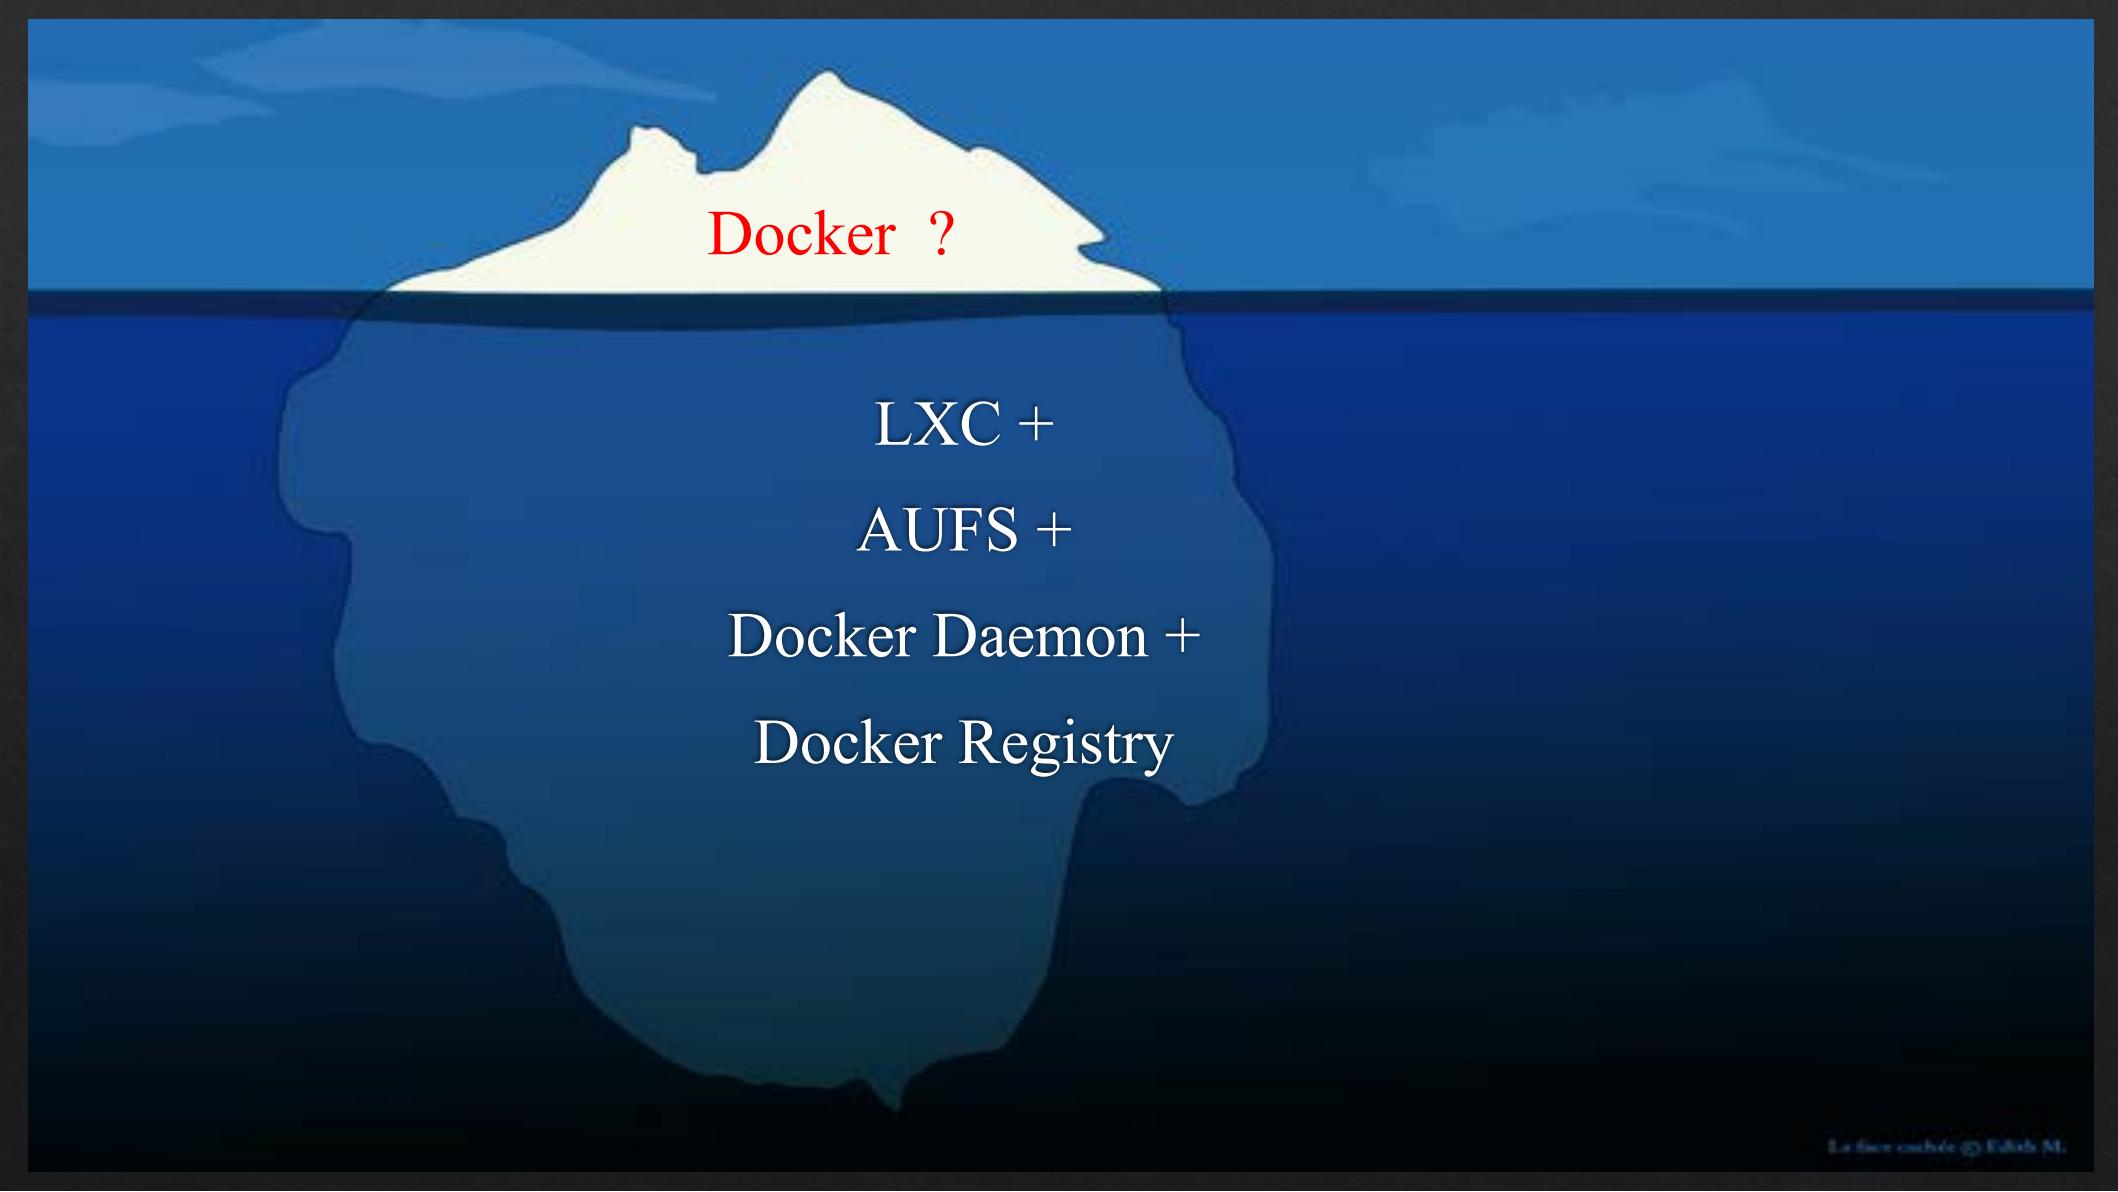


Docker glossary

Layers

- Images are based on a parent
- The layers stack on top
- Files in base layers are shared between Images
- Each commit creates a layer
- Base image has no parent



A large iceberg is shown floating in a dark blue sea under a light blue sky. The visible part above the water's surface is white and labeled "Docker ?" in red. Below the water's surface, the submerged portion of the iceberg is dark blue and contains the text "LXC +", "AUFS +", "Docker Daemon +", and "Docker Registry" stacked vertically.

Docker ?

LXC +

AUFS +

Docker Daemon +

Docker Registry

Containers are a userspace concept that takes advantage of several Kernel Subsystems

Key elements of Linux Containers

Process Isolation

Resource
Management

Security

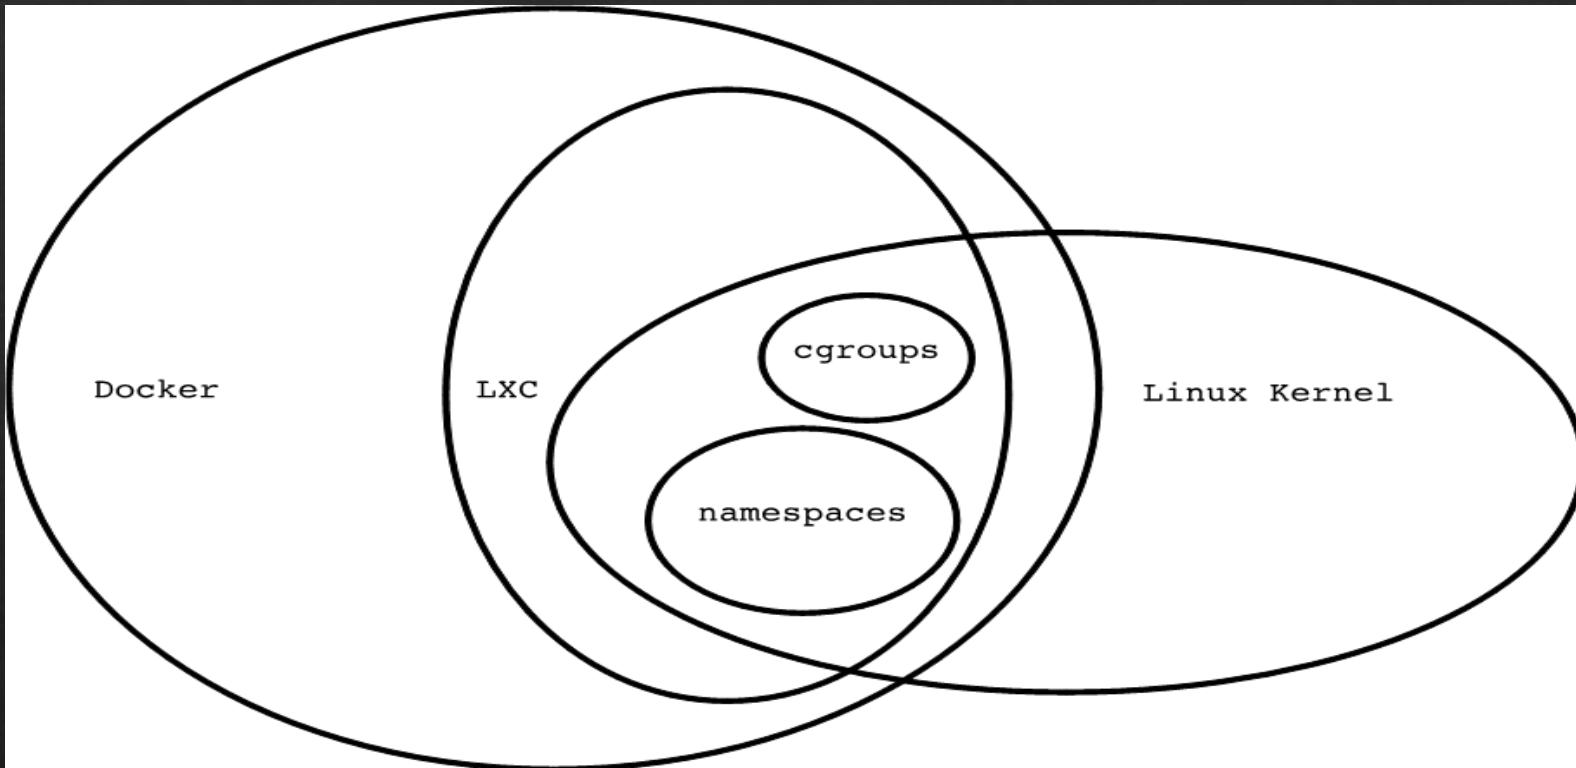
Management

How does this work

- cgroups
 - Group trees of processes
 - Allows control of resources for the group
- Namespaces
 - Isolate processes from host
 - Network, filesystem, pids, etc
- AuFS/DeviceMapper/Btrfs
 - CoW snapshotting of filesystem images



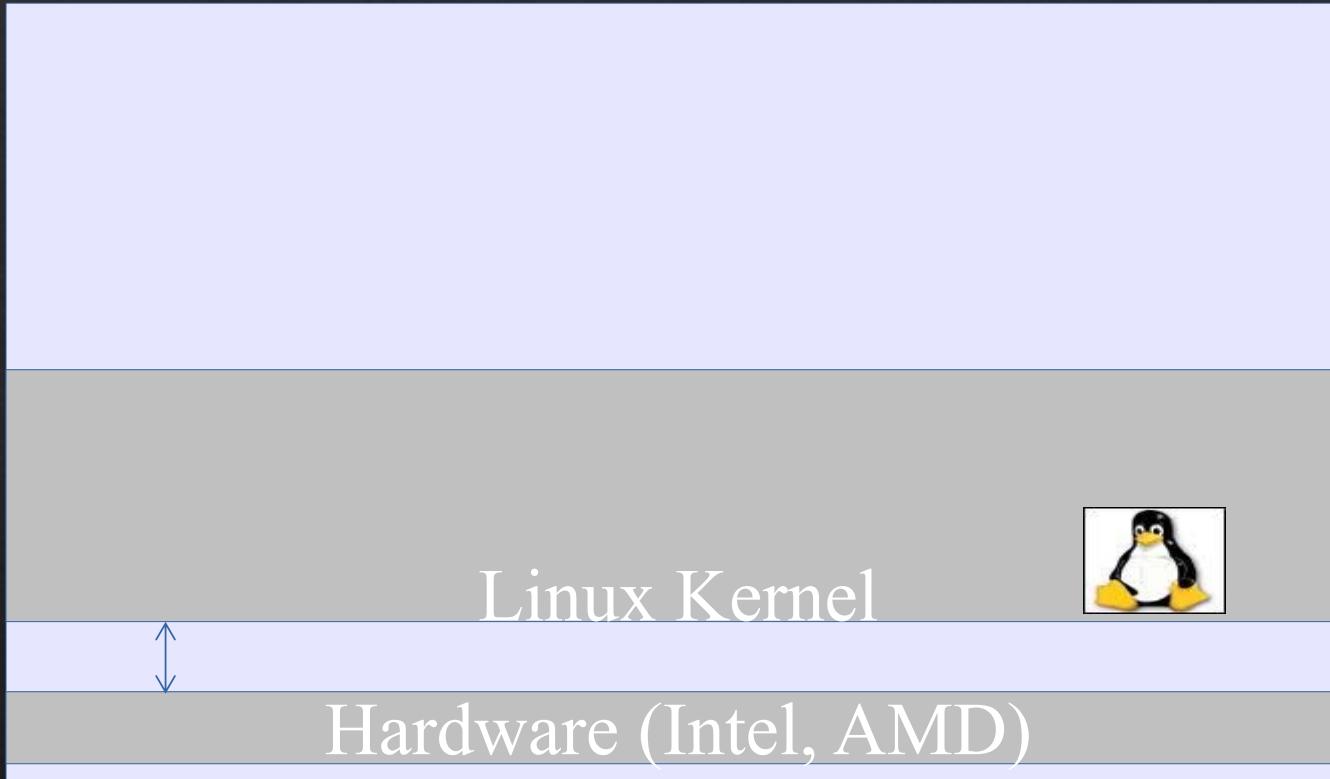
How does it work?



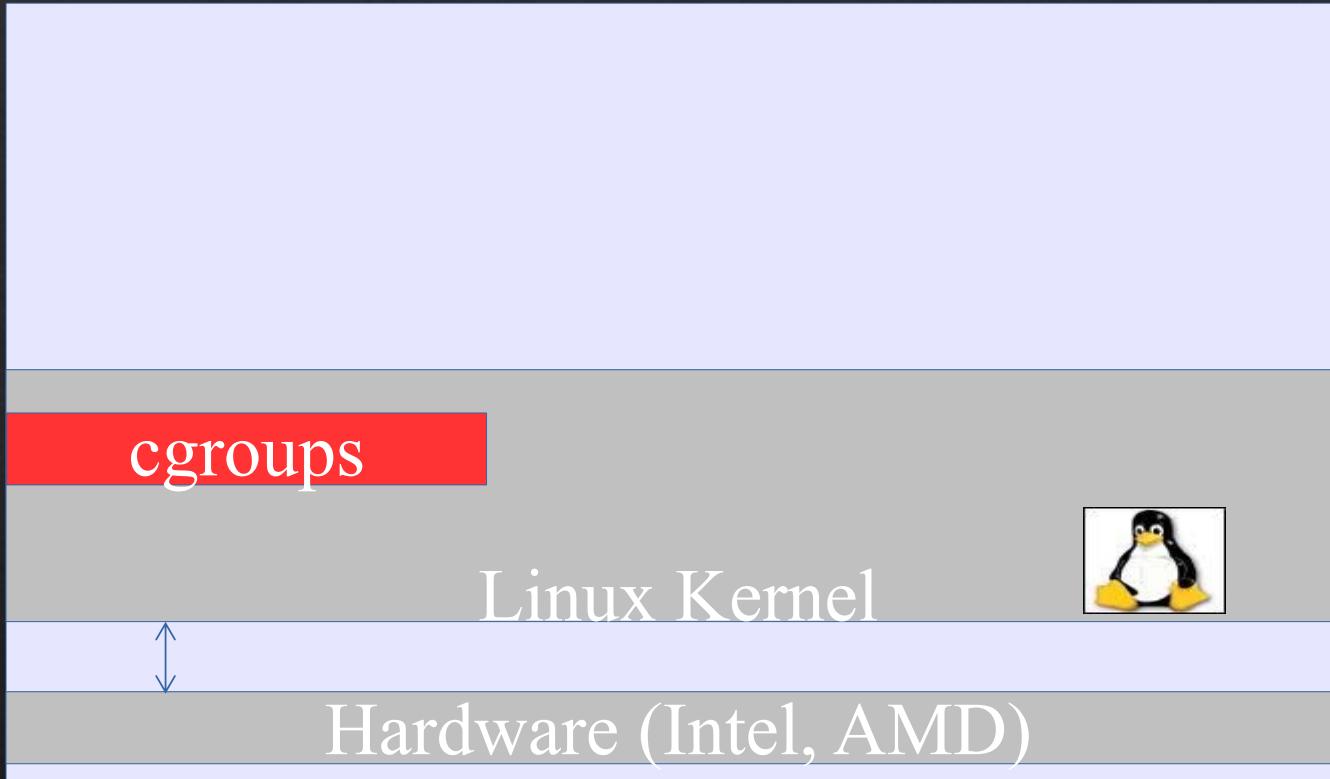
LXC (Linux Containers)

- Capability is already built into the kernel
- Utilizes cgroups (control groups) to limit, account and isolate resource usage (CPU, memory, disk I/O, etc.).
- **Stuff runs isolated from the rest of the host OS**

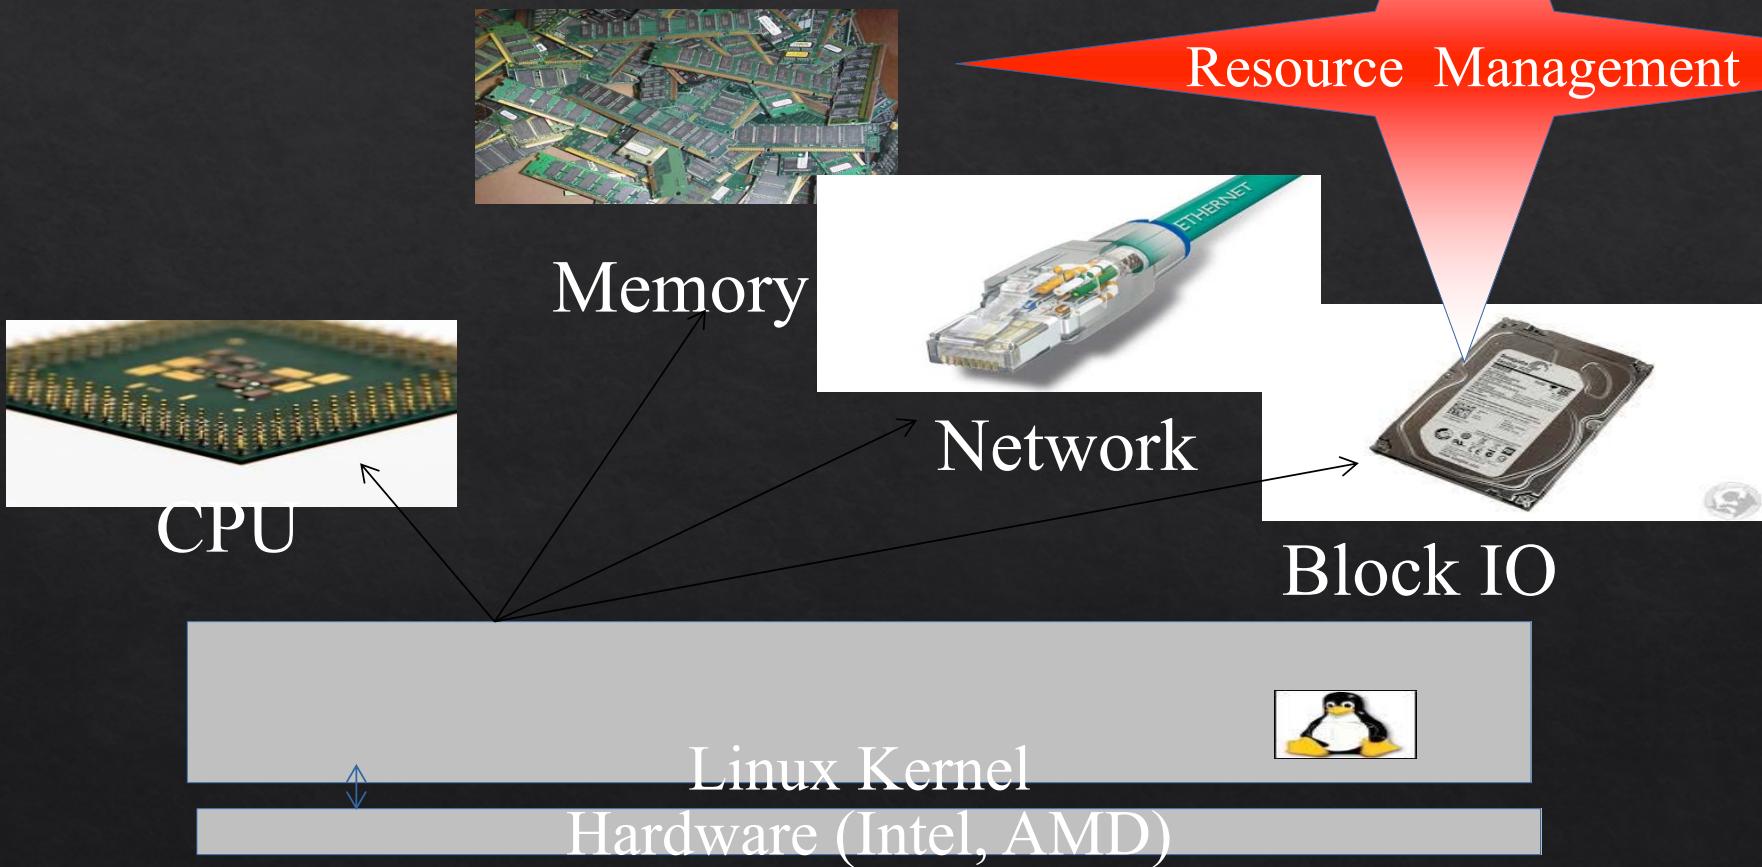
Linux Container Architecture



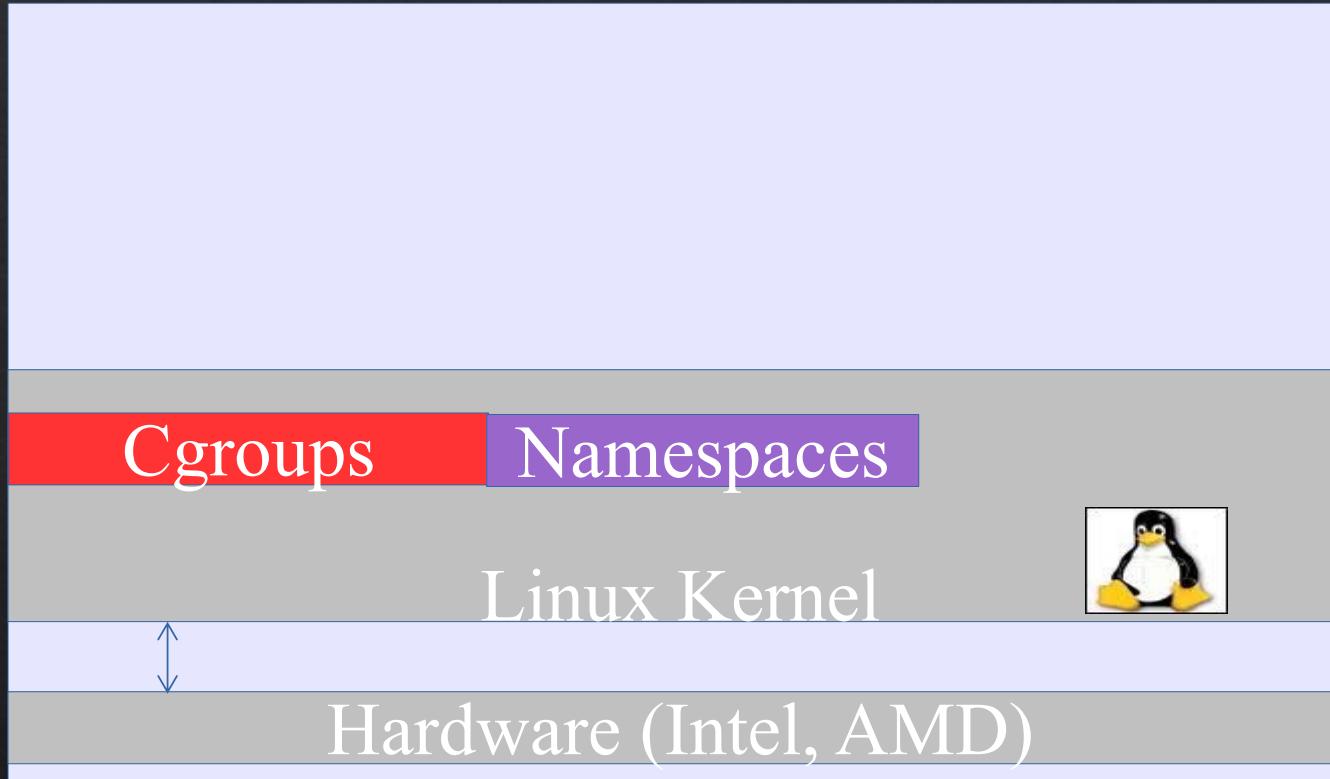
Linux Container Architecture



Cgroups

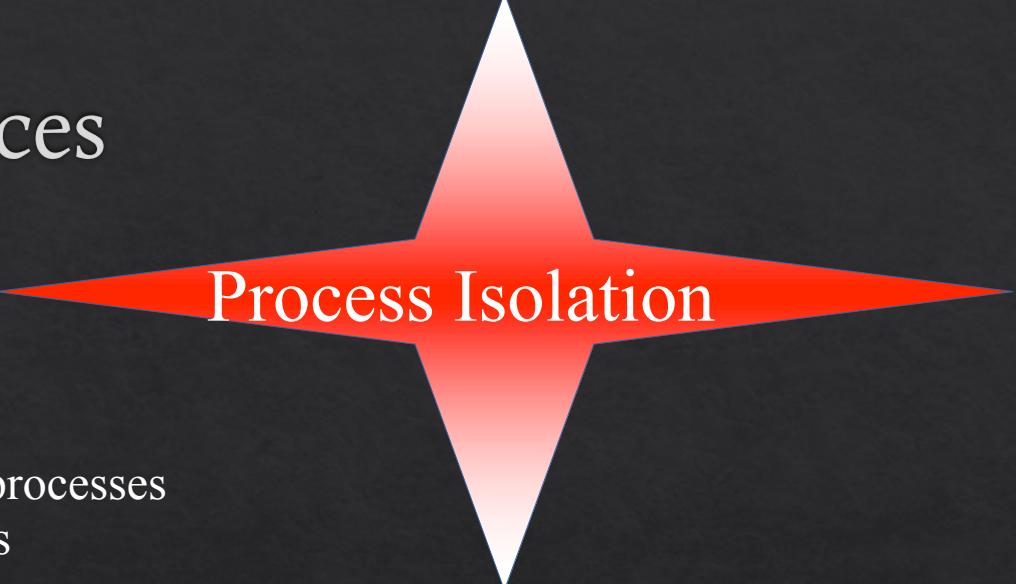


Linux Container Architecture



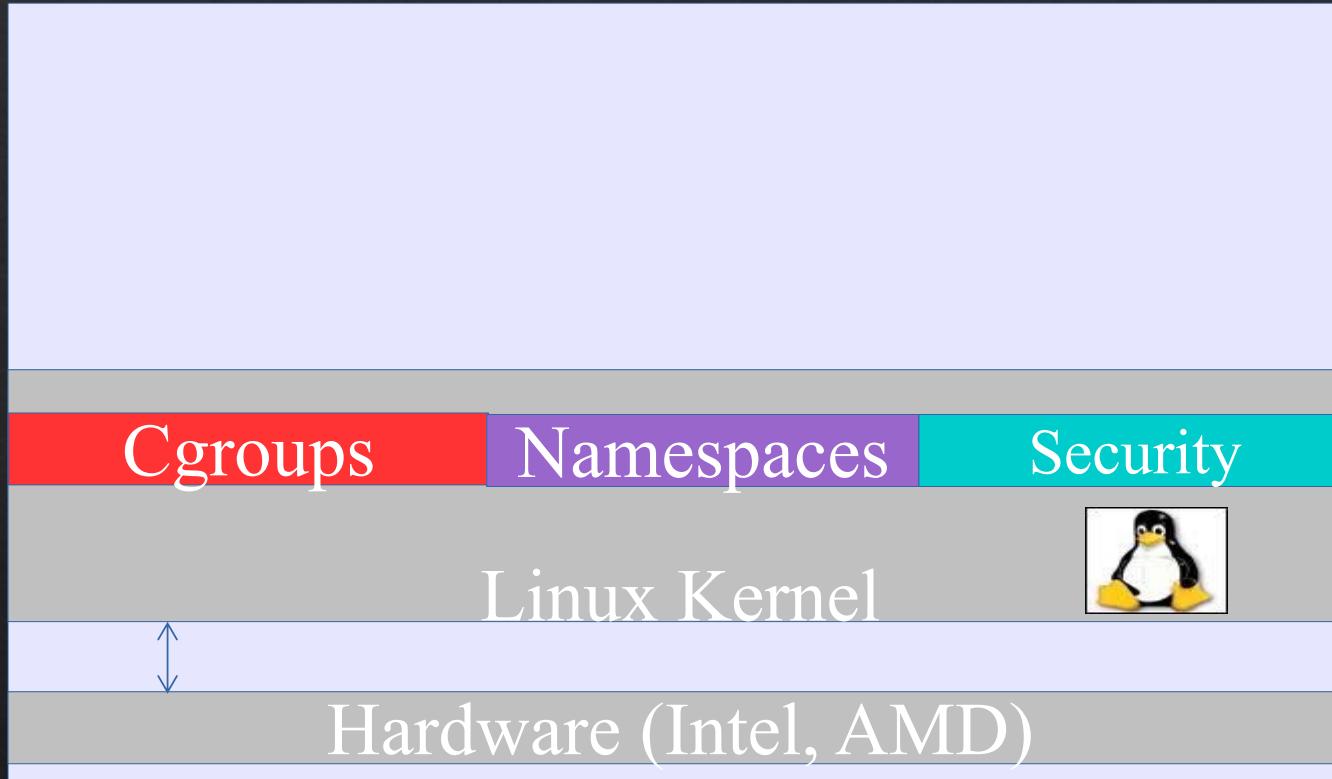
Namespaces

- Isolate processes
- Create a new environment with a subset of the resources
 - Once set up, namespaces are transparent for processes
 - Can be used in custom and complex scenarios
 - Supported Namespaces
 - ipc, pid, mnt, net, uts
- Future Red Hat Enterprise Linux 7: user

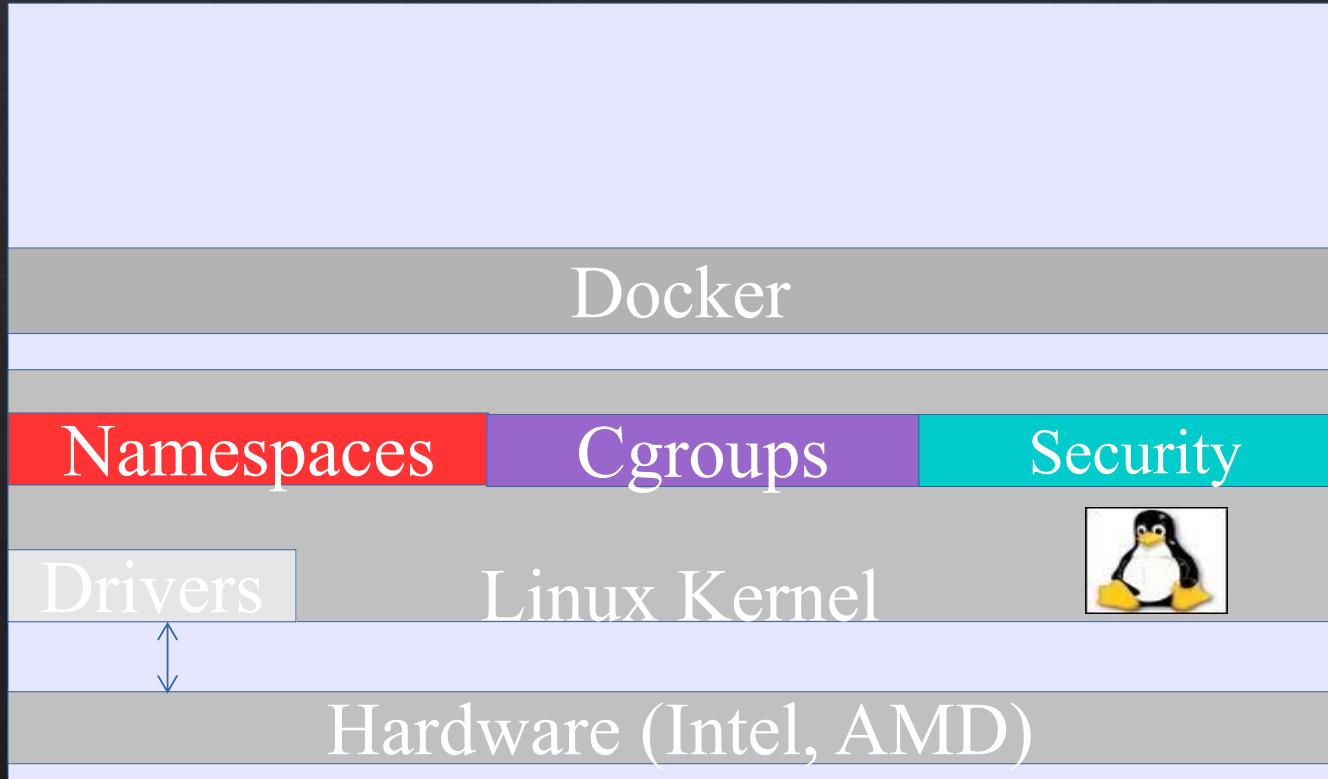


Process Isolation

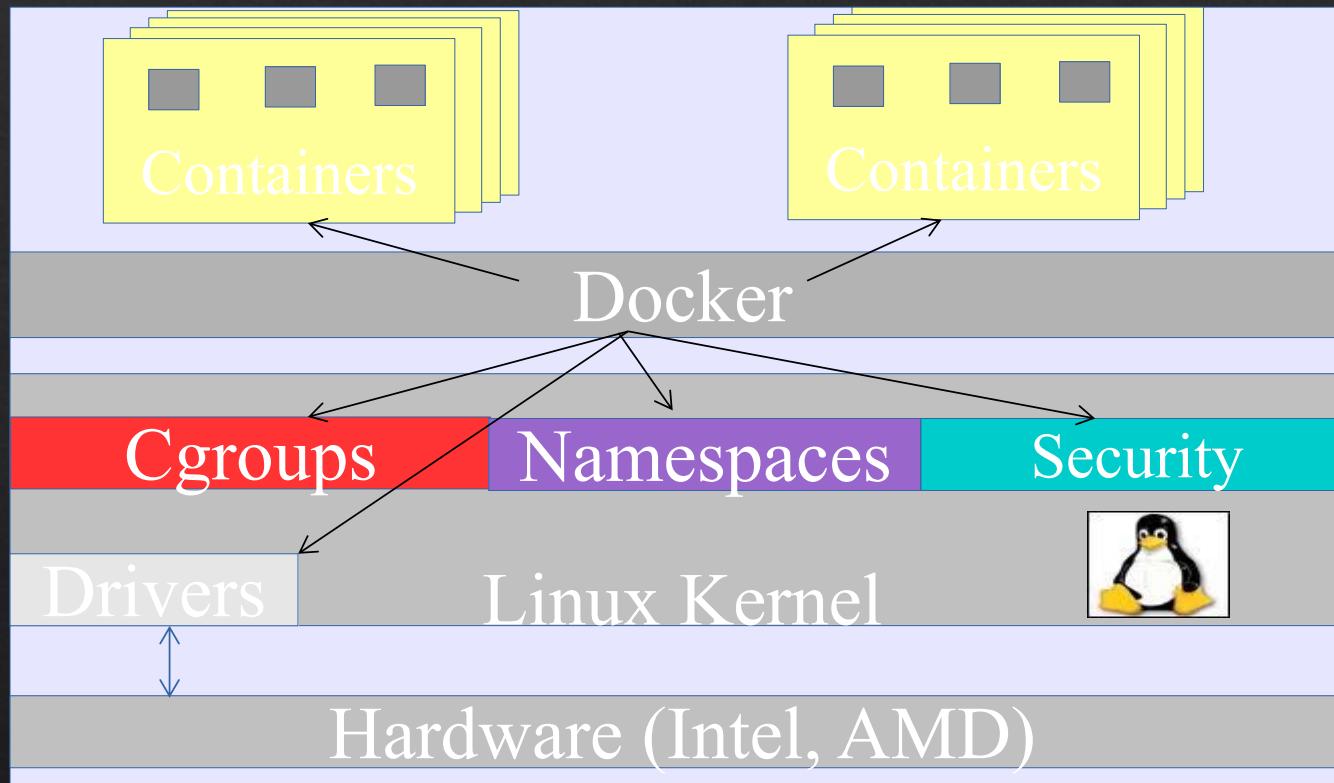
Linux Container Architecture



Linux Container Architecture

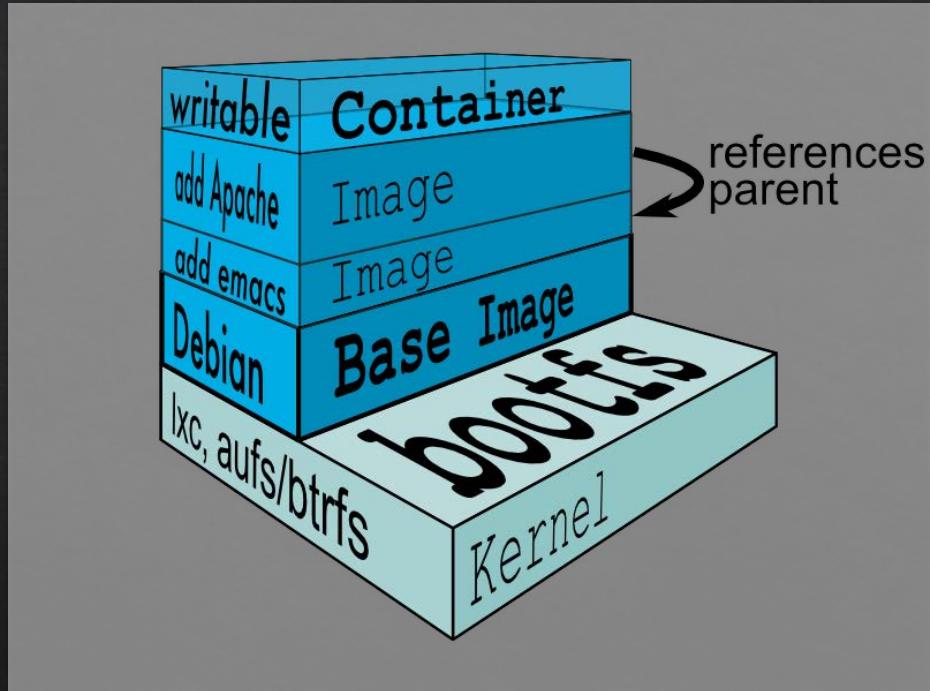


Linux Container Architecture



AUFS (AnotherUnionFS)

- Union file systems
 - Allows several file-systems to be mounted at one time, appearing to be one file-system.
- Docker uses it to create a layered file system.



Pretty picture from <http://docs.docker.io/en/latest/terms/layer/>

Docker Networking

- Docker uses a standard network bridge called docker0
- docker0 gets assigned an unused IP range
- Containers get bonded to docker0
- docker0's IP is the gateway for containers

Docker Daemon

- Software layer that allow for easy creation and management of Docker containers (glorified LXC instances).

Example Docker Work Flow (CentOS + SSH)

```
docker pull centos
```

```
docker run -t centos yum -y install openssh-server
```

```
docker commit <image id> centos/ssh
```

```
docker run -d -p 2222:22 -t centos/22 /usr/sbin/sshd
```

Host/Container Communication

- Ports can be exposed and mapped to the host with “docker -p”
 - docker run -p 80:80 -t <image> <cmd>

Container to Container Communication

- Handled with container linking
 - docker run -d -p 6666 --name parent <image>
 - docker run --link parent:child -i -t <image>
- Linked container will have environment variables like:
 - CHILD_PORT_6666_TCP_ADDR=172.17.0.2
 - CHILD_PORT=tcp://172.17.0.2

Docker Registry

- Software that runs the docker index at <https://hub.docker.io/>
- Base images for things like CentOS, Ubuntu, Fedora, etc are pulled from the public registry.
- Local registries can be created and submitted to.

Docker Files

- Allow for easy recreation of an application anywhere Docker can be ran. Example:

```
# use the ubuntu base image provided by dotCloud
FROM ubuntu

# make sure the package repository is up to date
RUN echo "deb http://archive.ubuntu.com/ubuntu precise main universe" > /etc/apt/sources.list
RUN apt-get update

# install memcached
RUN apt-get install -y memcached
```

Uses for Docker

- Testing on multiple versions of multiple distros without the pitfalls of standard virtualization.
- Running newer or older versions of applications not in your host OSes repos.
- Creating an applications that can be easily rebuilt and reused on any distro, anywhere.

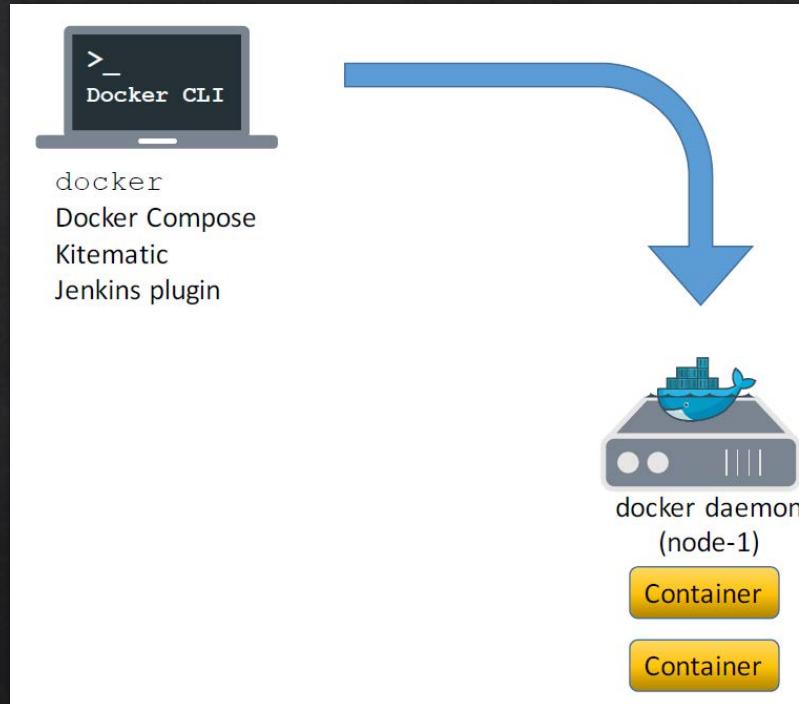
Docker Swarm

Running docker on a cloud

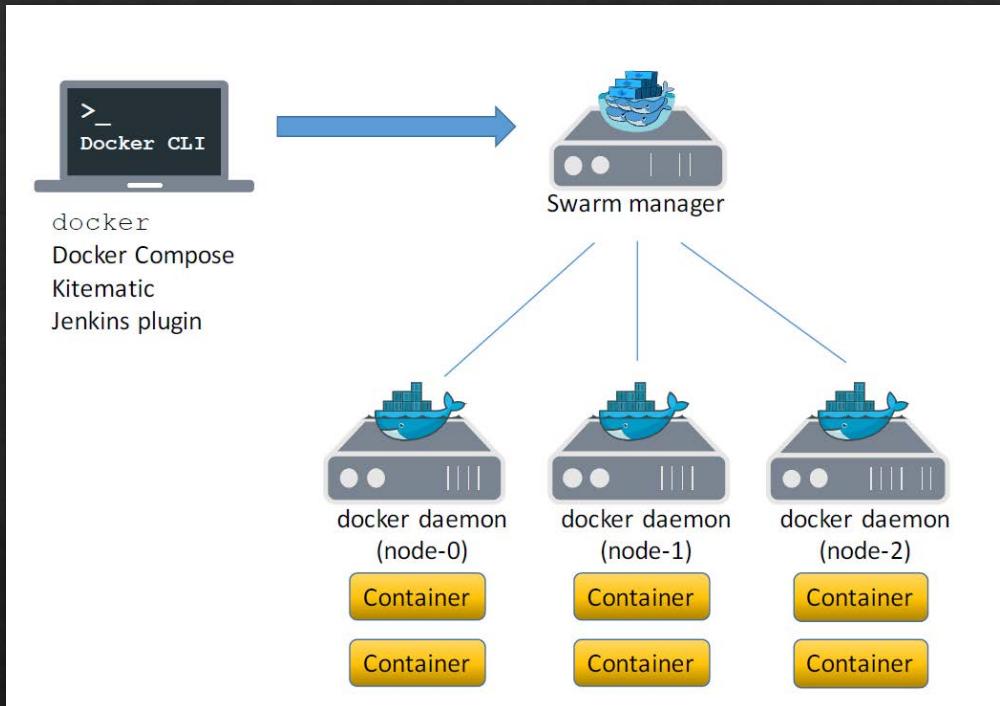
What is Swarm?

- ◆ Swarm turns multiple Docker hosts into a single, virtual Docker host.

Classical docker usage



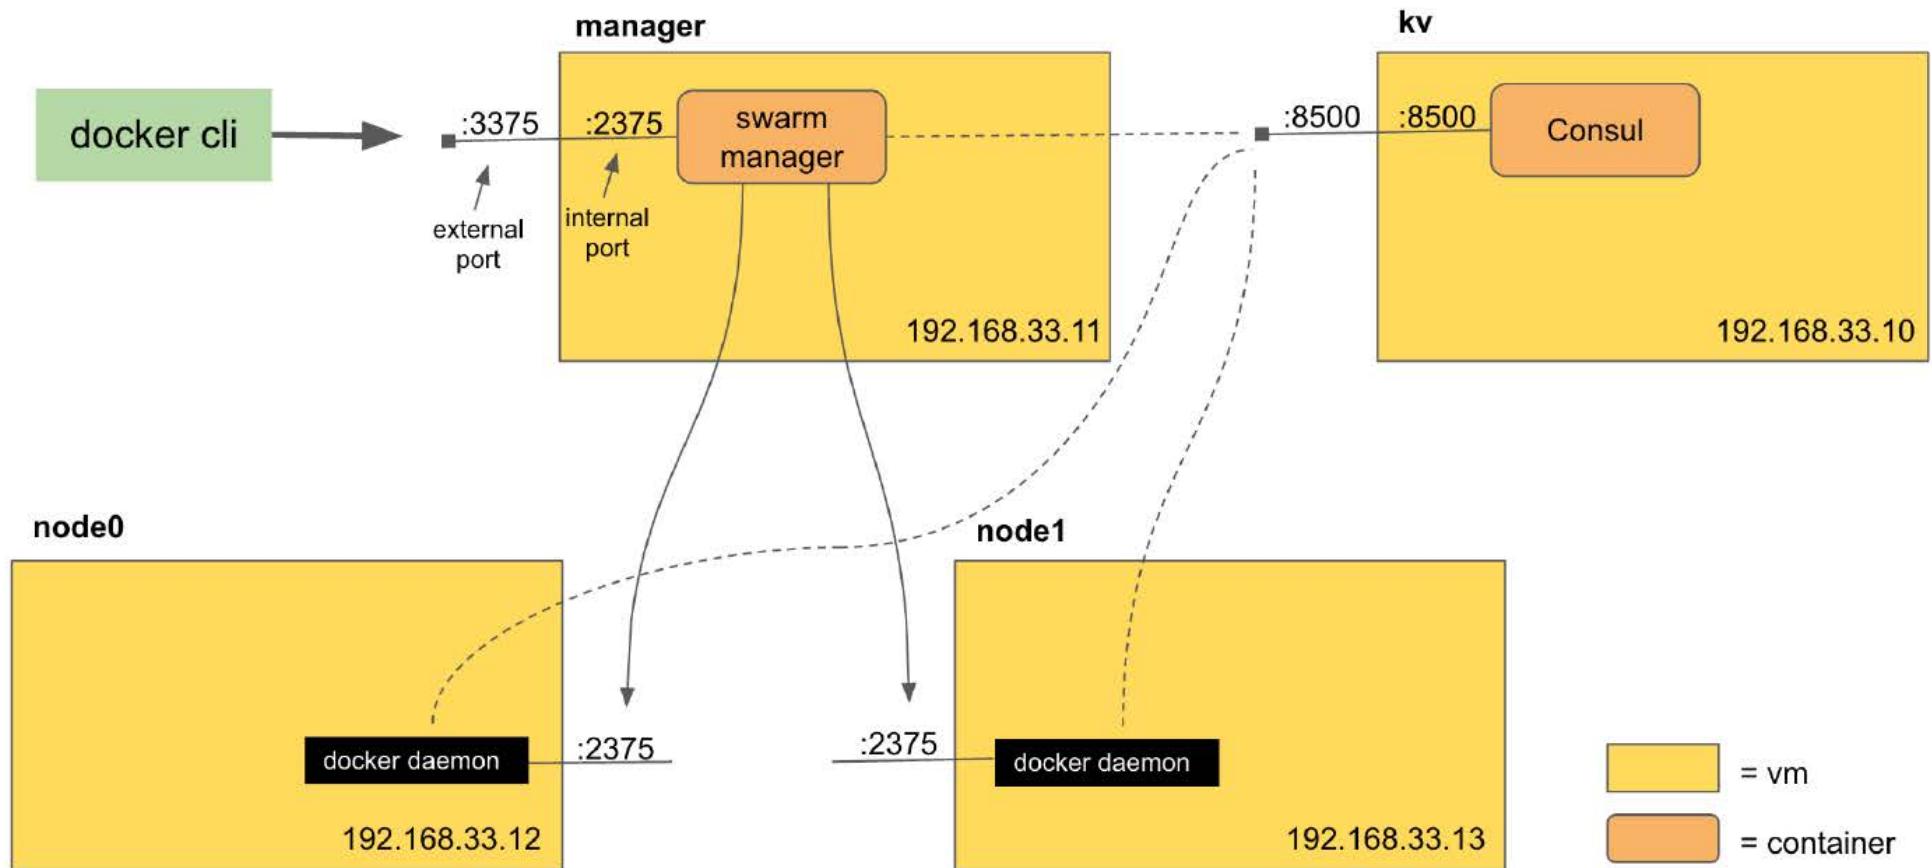
... with Swarm



Swarm features

- ❖ Scheduling and rescheduling on failure
- ❖ Labels, affinity and constraints to guide the placement of containers

Swarm Cluster



3 steps to run your swarm cluster

- ❖ 1/ Create a KV (for node discovery and lib networkIP allocation)
- ❖ 2/ Run your Swarm manager (on a single node)
- ❖ 3/ restart your Docker daemon with swarm parameters (on all your worker nodes)

Docker Compose

Multi-container apps are a hassle.

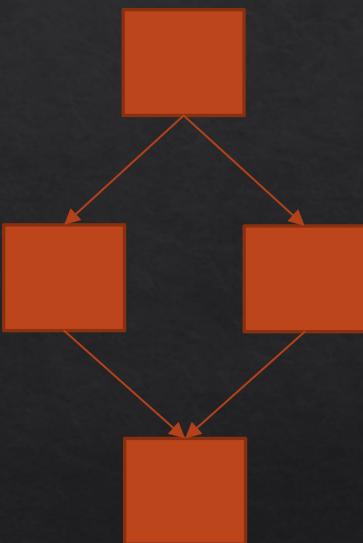
- ❖ Build images from Dockerfiles
- ❖ Pull images from the Hub or a private registry
- ❖ Configure and create containers
- ❖ Start and stop containers
- ❖ Stream their logs



Multi-container apps are a hassle.

```
$ docker pull redis:latest  
$ docker build -t web .  
$ docker run -d --name=db redis:latest redis-server --appendonly yes  
$ docker run -d --name=web --link db:db -p 5000:5000 -v `pwd`:/code web python app.py
```

Multi-container apps are a hassle.



\$ docker pull ...

\$ docker pull ...

\$ docker build ...

\$ docker build ...

\$ docker run ...

\$ docker run ...

\$ docker run ...

\$ docker run ...

Docker Compose:

Get an app running in one command.



Native Docker API

- ❖ The same API used by **other Docker commands**.
- ❖ The same API used by **other Docker tools**.
- ❖ The same API **exposed by Docker Swarm**.

- ❖ **Multi-container apps on multi-host clusters.**

Built into the Docker client

- ❖ New command: docker up
- ❖ Enhanced commands: docker build, pull, run, start, stop, kill, rm...

A small toy example of docker compose file

Create a file called *.yml in your project directory :

This Compose file defines two services, web and redis. The web service:

- ❖ Uses an image that's built from the Dockerfile in the current directory.
- ❖ Forwards the exposed port 5000 on the container to port 5000 on the host machine.
- ❖ Mounts the project directory on the host to /code inside the container, allowing you to modify the code without having to rebuild the image.

```
version: '2'
services:
  web:
    build: .
    ports:
      - "5000:5000"
    volumes:
      - .:/code
  redis:
    image: "redis:alpine"
```

The redis service uses a public Redis image pulled from the Docker Hub registry.



Summary

Advantages:

- Portable configuration
- Reuse images other people have built
- Lightweight, fast
- Easy to scale up
- "Build once run everywhere"

Disadvantages:

- Fast moving

Thank You



<https://www.docker.io/>